

Speech Diarization using Deep Neural Networks and Machine Learning

Clint Ferrin, Madi Mickelsen, Chui Vanfleet, Daniel Mortensen

December 13, 2017

Contents

1	Introduction	3
1.1	Problem	3
2	Requirements	3
2.1	Algorithm Requirements	3
2.2	Report Requirements	3
3	Methods	4
3.1	Fully Connected Neural Network	4
3.2	Blind Source Separation	6
3.3	Transfer Function Estimation	6
3.3.1	Methodology	7
3.3.2	System Parameters	8
4	Results	9
5	Conclusion	11
A	Code Listings	14

A.1	main.py	14
A.2	DiarizationBaseClass.py	18
A.3	TFLearnNN.py	20
A.4	plot.py	22
A.5	textgrid.py	25

List of Figures

1	Block Diagram of Neural Network	5
2	Fully Connected Neural Network	5
3	Overall System Diagram	7
4	Non-Zero Error System	7
5	Zero Error System	8
6	Neural Network System Design	8
7	Cross Correlation Data	9
8	TensorBoard Loss Training Plot	9
9	TensorBoard Accuracy Training Plot	10
10	Example of Poor Diarization	10
11	Results of Network on Untrained Data	11

1 Introduction

“Audio diarization is the process of annotating an input audio channel with information that attributes (possibly overlapping) temporal regions of signal energy to their specific sources” [1]. This report represents an effort to solve speaker diarization problem using deep learning. Many other methods of speech diarization rely heavily on signal processing techniques, but fewer attempt to use neural networks [2, 3, 4].

Dr. Stephanie A. Borrie from the Human Interaction Lab at Utah State University has provided a number of audio files from her research. The audio files contain a recording of two individuals conversing. Each individual speaks into their own microphone. But, crosstalk from each speaker is picked up on the other microphone.

Along with the audio files, Dr. Borrie has provided TextGrid files which contain the diarization annotations. These TextGrid files were created by Dr. Borrie’s students and are to be used as training and testing data for a deep learning algorithm.

1.1 Problem

Using the audio and TextGrid files provided by Dr. Borrie, a deep learning algorithm must be developed to automatically perform speaker diarization

2 Requirements

2.1 Algorithm Requirements

1. The project must use a neural network at some point. If additional processing steps are used, such as in preparation of the data, or organizing the output response, that is also fine.
2. The training code should provide some indication of how the training is progressing.
3. The code should write a file indicating the diarization.

2.2 Report Requirements

1. Clearly describe the architecture used in processing.
2. Clearly describe the neural architecture, such as number of layers, neurons in each layer, type of neuron, number of weights, stride, etc.
3. Clearly describe what the inputs are.
4. Clearly describe how many outputs there are, and what the output layer function is (sigmoidal, softmax, linear, etc.).

5. Clearly describe the cost function the neuron is training against, and the kind of optimization used.
6. Clearly describe any other processing steps used in conjunction with the neural net.
7. Describe how to train the neural net — how it is presented input data, how the output data is presented.
8. Present data, preferably learning curves, demonstrating the neural network is learning. Indicate how many training iterations used, and why.
9. Present results quantifying how well the diarizer works. Ideally, this would be compared to the original training data. Provide a description of how the score is derived.
10. Provide a description of what was learned.
11. Provide a printout of the code.

3 Methods

Several different diarization methods were explored to meet the requirements discussed in subsection 2.1. This section will discuss, in detail, the several diarization methods created, presented in no particular order.

3.1 Fully Connected Neural Network

The fully connected version of the Neural Network can be seen in Figure 1. The Network has 5 layers, it receives a total of 3400 inputs (1600 inputs from channel 1, 1600 inputs from channel 2, 100 inputs from the `fft` of channel 1, and 100 inputs from the `fft` of channel 2), and has four possible outputs (00-no speech, 01-speaker one, 10-speaker two, and 11-both speakers).

The input layer is stacked horizontally with channel 1 sampled at 8000 Hertz for a total length of 0.2 seconds, channel 2 sampled at 8000 Hertz, and finally both channels are followed by their corresponding `fft` results with a size of 100 each. This yields a total input size of 3600; each new sample of training or testing data is stacked horizontally on top off each input sample to be fed into the network—the stacking of input data in matrix form allows TensorFlow to speed up the training process.

In order from left to right, each layer has 300, 200, 200, and 500 neurons. All of the hidden layers utilize ReLU activation functions, and the densely connected layers 2 through 4 have a dropout rate of 0.8. The output layer uses the softmax activation function because there are 4 different classes that the loss function needs to categorize.

To train the network, a file path is given to the main function. The number of `.wav` files in the audio files directory must match the number of `.TextGrid` files in the labels

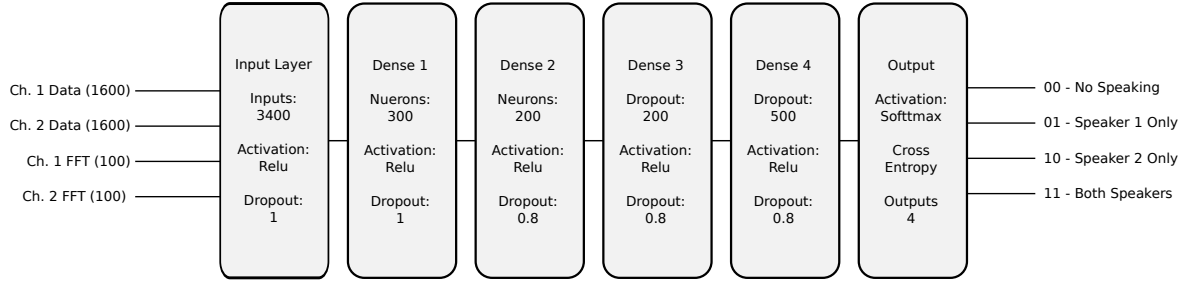


Figure 1: Block Diagram of Neural Network

directory. The label **numpy** array is created by sampling the **TextGrid** file to create an output list that corresponds to every F_s (samples per second) in the audio file.

The cost function utilized in the network design is cross-entropy, and the output of the network is compared to a one-hot-encoded **Y** vector that contains the actual class that was marked by humans. The optimization function that performed best with this setup was the **AdamOptimizer** that is included in the TensorFlow parameters. For more details reference the TensorBoard printout in Figure 2

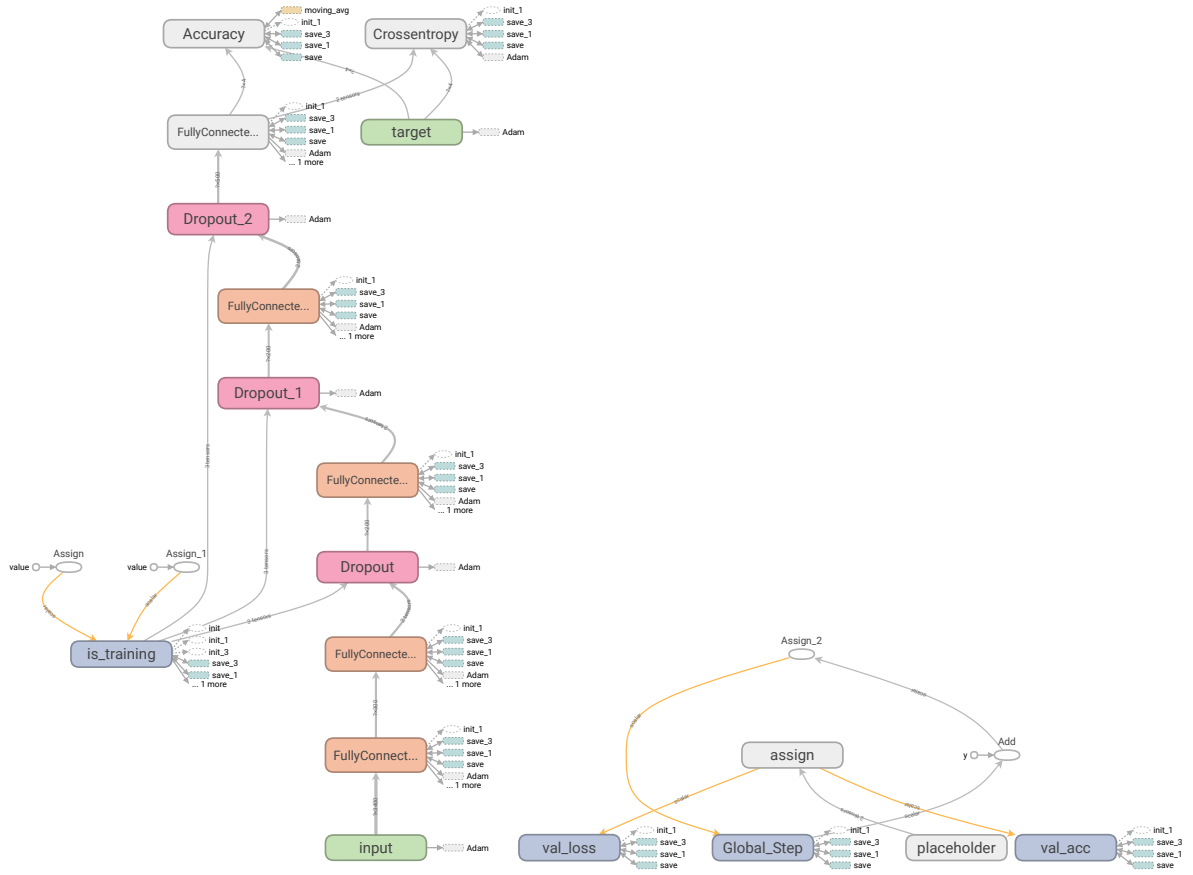


Figure 2: Fully Connected Neural Network

While implementing the network, different numbers of neurons, dropout rates, and

network types were experimented with. The 1D convolutional network did not produce better results, and if the dropout rate was too high, the loss function would not converge. The parameters were changed and re-run to identify the best network system.

Introducing the FFT significantly improved the accuracy of the network. Before the FFT, the network had a very difficult time identifying silence, but the added information from the FFT improved the overall performance. For the results of this method, see Section 4

3.2 Blind Source Separation

The original design included a preprocessing portion using blind source separation (BSS). The goal of BSS, also known as independent component analysis, is to determine a matrix which forces the elements of the output to be statistically independent. The matrix is determined using an iterative process of maximum likelihood independent component analysis. When multiplied by the output, each element should be separated from the others.

For the purpose of this project, BSS would be used to separate the two speakers when speech overlap occurs, removing crosstalk. The BSS program was developed for another class using MATLAB. Before implementing BSS into the design fully (converting the MATLAB program into python), segments of an audio file with overlapped speech were run through the BSS program developed in MATLAB. This would test the BSS method to determine if it would add, or detract from the system.

Initially, the program would not separate the signals whatsoever. The audio appeared to be of worse quality, with more crosstalk rather than less. More iterations were run, the step size was changed, and more segments were concatenated to see if more data worked, but none of these possible *improvements* did not appear to change anything.

The reason BSS did not work for this instance could be due to the delay between microphones. For example, speaker one speaks directly into microphone one, whereas speaker two speaks directly into microphone two. The distance between the speakers would cause a delay to occur between speaking into one microphone and the crosstalk on the second. The BSS algorithm used in the program does not account for the delay. As a result this method was not incorporated into the final design of the system.

3.3 Transfer Function Estimation

Based on previous results, the following is a writeup of the next step that is predicted to create better results than the previous. In past attempts, windows of data were inserted into a neural network, along with their corresponding FFTs, to determine where changes in speaking would occur. This was complicated by microphone recording speech signals from several participants at once. This section goes over a method designed to show which microphone is being used during a given time interval.

3.3.1 Methodology

This method operates under the assumption that the input to a single microphone consists of the input from the nearest speaker and the output from the second speaker after it has been run through a transfer function $H(Z)$ as seen in Figure 3. This portion




Figure 3: Overall System Diagram

of the speech diarization would be purposed towards detecting who is speaking: person one, person two, or both. For intents of the following setup, it is assumed that one person is speaking at a time. This algorithm requires two channels, one for microphone one, and the other for microphone two. In this setup, if only one person is speaking, then one signal will represent the most accurate data from the source, and the other will represent the same data after it has passed through a system and been modified by the system's transfer function.

The proposed solution is to develop a neural network that would simulate the transfer function that operates on the non-dominant dataset. Once this is accomplished, channel A is run through the simulated transfer function and subtracted from the other channel. If the output is minimal, then the channel that was operated on matches its user. In other words, when person one's data is run through the transfer function, and subtracted from person two's data, if the output is close to zero, then person two was not talking, and person one provided input to the microphone and vice versa. Figure 4 demonstrates a system view if the input signal was pre-transformed. Figure 5 shows the




Figure 4: Non-Zero Error System

system view if the input signal is not passed through the transfer function, for example, it is person one's signal when person one was speaking. Channel two should then be run on input data. The results from both systems should be compared. If non-zero values




Figure 5: Zero Error System

are detected after running both channels through the transfer function, then both were talking. An example of the proposed system is shown in Figure 6. Both the original sig-

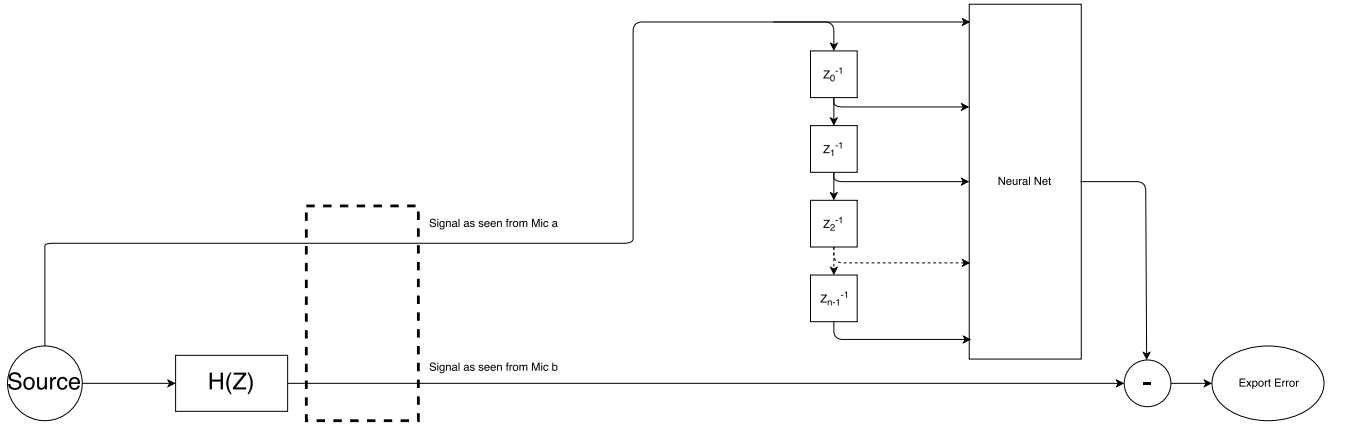


Figure 6: Neural Network System Design

nal is provided as input ,and the time delayed versions of the same signal. It is assumed when the signal undergoes a transformation, the transformed signal is consists of linear combinations of time delayed source samples.

3.3.2 System Parameters

The neural network will be comprised of a semi-connected layer where each neuron has access to a time delayed sample of the input data. The data is arranged as a Toeplitz matrix, A , with each row representing the input data of each neuron. The output of the network has equivalent dimensions as the output of the transfer function. There should be between 0 - 500 sample delays based on cross correlation, as shown in Figure 7. Stronger correlations between signals are seen from as far as $\tau = 500$ samples out.

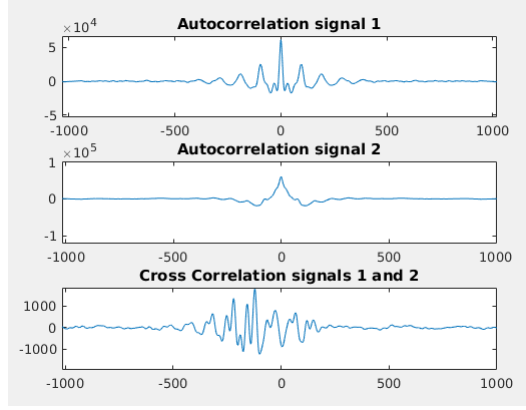


Figure 7: Cross Correlation Data

4 Results

Using the network described in Section 3.1, the network correctly labeled a file it had not seen with 75% accuracy. The training produced the following loss function graph after 25 epochs. 25 epochs was chosen, with batch sizes of 100 because larger batches or epochs cause the network start to decline in accuracy due to over-fitting. The loss plot in Figure 8 converges to a value; with more epochs, little progress is made in reducing the loss.




Figure 8: TensorBoard Loss Training Plot

The output of the accuracy plot can be seen in Figure 9. The accuracy was obtained by splitting the data into random training and testing sets of length 0.2 seconds. To create the training and testing data sets, all of the data was placed into a long array of data and shuffled in the same way as the labels. The data was down-sampled to a new rate of 8000 Hz, and the labels were sampled at 5 Hz. After the shuffling, the data was split into the proportion indicated by the variable `percent_train`.

As a result, the accuracy represents categorizing data that has not been trained by

the Neural Network.




Figure 9: TensorBoard Accuracy Training Plot

It should be noted that the human markings were incorrect in many of the text grids. For example, the file `HS_D09.wav` seen in Figure 10 shows the first 10 seconds of a file that has been poorly marked. When listening to the file with both channels, it is obvious that no one is speaking for the majority of the 10 seconds. In that audio portions that people are speaking, the markings are incorrect. If the TextGrid files were more accurately labeled, the output of the programs would be more accurate as well.




Figure 10: Example of Poor Diarization

To test the effectiveness of the training, the network was given a complete file of

4:00 minutes, and told to predict the labels. Despite the poorly marked training data, the network achieves 75% accuracy on an entire data file that the network had not seen previously. To test the accuracy, the diarization state was predicted every 0.2 seconds and compared to the markings on each tier of the TextGrid files. As a visual representation, the output after the network predicted the labels of the file was graphed. Figure 11 shows all of the necessary information to view how the network performed.




Figure 11: Results of Network on Untrained Data

The solid fill colors represent where the human labeled the speaking of both people. The areas with no color represents when neither person one or person two are speaking, blue fill represents when person one is speaking, red fill represents when person two is speaking, and green fill represents when both people are speaking at the same time.

The solid lines at the bottom of the graph are the predictions that the network produces. The predictions of the network line up closely to the human markings. In many cases, the neural network identifies small sounds that people make that are not recorded as speech in the TextGrid file (such as brief laughs). This can be filtered out at a future date to improve accuracy.

5 Conclusion

The fully connected neural network with the fft inputs was used to label audio and performed with 75% accuracy. The fully connected neural network was chosen as our

final design because it yielded the best results. TensorFlow was a challenge to use by itself and it was difficult to quickly change network design, so using APIs like TFLearn and Keras made the process much easier.

Alternate methods to the speech diarization problem were discussed once the initial design was implemented and tested. These would be implemented in future iterations of this project to determine if they improved the accuracy.

The first alternate approach would use the Neural Network for the BSS. The NN should be able to incorporate the delay, making it more robust than the BSS algorithm used initially. This would allow easy separation of the two channels. After filtering out the noise, the speech on each channel would be easily detectable using the power of the signal.

Section 3.3 discussed another method of approaching the problem, estimating the transfer function. The transfer function estimation method would help to identify when each microphone is being used but may prove difficult to implement depending on the different recording environments.

References

- [1] S. E. Tranter and D. A. Reynolds. “An overview of automatic speaker diarization systems”. In: *IEEE Transactions on Audio, Speech, and Language Processing* 14.5 (Sept. 2006), pp. 1557–1565. ISSN: 1558-7916. DOI: 10.1109/TASL.2006.878256.
- [2] K. Boakye et al. “Overlapped speech detection for improved speaker diarization in multiparty meetings”. In: *2008 IEEE International Conference on Acoustics, Speech and Signal Processing*. Mar. 2008, pp. 4353–4356. DOI: 10.1109/ICASSP.2008.4518619.
- [3] X. Anguera et al. “Speaker Diarization: A Review of Recent Research”. In: *IEEE Transactions on Audio, Speech, and Language Processing* 20.2 (Feb. 2012), pp. 356–370. ISSN: 1558-7916. DOI: 10.1109/TASL.2011.2125954.
- [4] C. Barras et al. “Multistage speaker diarization of broadcast news”. In: *IEEE Transactions on Audio, Speech, and Language Processing* 14.5 (Sept. 2006), pp. 1505–1512. ISSN: 1558-7916. DOI: 10.1109/TASL.2006.878261.

A Code Listings

A.1 main.py

```
1 import tensorflow as tf
2
3 from scipy.io import wavfile
4 from textgrid import TextGrid
5 from diarization.methods.TfLearnNN import TfLearnNN
6 from os import walk
7 from os.path import splitext
8 from numpy import array, vstack, hstack, zeros, ceil
9 from numpy.random import shuffle, randint
10 import time
11
12 audio_sample_rate = 8000
13
14 # Parameters to be set in some more elegant way later on.
15 audio_files_path = '../media/Full_Test_Files_8000/'
16 text_grids_path = '../media/Full_Test_TextGrids/'
17 load_all_to_memory = True
18 window_size = int(audio_sample_rate * 0.2)
19 num_classes = 4
20 num_channels = 2
21 # training_flag = True
22 training_flag = False
23 epochs = 50
24 batch_size = 100
25 percent_train = .8
26 params = {
27     'num_channels' : 2,
28     'window_size' : window_size,
29     'num_classes' : 4,
30     'nfft' : 100
31 }
32 diarization_method = TfLearnNN(params)
33 load_path = '../media/tflearn_nn.save'
34
35 # run_audio_file_path = '../media/Full Test Files 8000/HS_D01.wav'
36 # text_grid_load_file_path = '../media/Full Test TextGrids/HS_D01.TextGrid'
37 # text_grids_save_file_path = '../media/TextGrid Output/'
38
39 # save_path = '../media/tflearn_nn_2.save'
40 save_path = None
41 # load_path = None
42 load_path = '../media/tflearn_nn.save'
43
44 # End parameters section
45
46
47 def get_file_list (directory):
48     """
49     Used to get all the file names in the directory specified .
50     :param directory: the path to the directory to get the file names from (ending in '/')
51     :return: a list of all files names in directory starting with directory
52     """
```

```

53 f = []
54 for (dirpath, dirnames, filenames) in walk(directory):
55     f = filenames
56     break
57 f.sort()
58 for i in range(len(f)):
59     f[i] = directory + f[i]
60 return f
61
62
63 def get_label(path):
64     """
65     Used to load a TextGrid from disk and return the array representation of it.
66     :param path: Path to the TextGrid
67     :return: Array representation of the TextGrid.
68     """
69     grid = TextGrid(name=path)
70     grid.read(path, Fs=int(audio_sample_rate / window_size))
71     return grid.FsArrayCombined
72
73
74 class Trainer:
75     def __init__(self):
76         self.data_paths = None
77         self.label_paths = None
78         self.training_data = None
79         self.training_labels = None
80         self.testing_data = None
81         self.testing_labels = None
82         self.imax = -1
83
84         self.data_paths = array([ get_file_list ( audio_files_path )])
85         self.label_paths = array([ get_file_list ( text_grids_path )])
86         if len( self.data_paths) != len( self.label_paths):
87             print('Error!_Data_and_labels_do_not_match_up!')
88             ValueError()
89
90         self.split_train_test ()
91
92         if load_all_to_memory:
93             print('Loading_Data...', end='', flush=True)
94             self.load_all_data_to_memory()
95             print('Done!')
96         else:
97             print("Training_is_going_to_take_a_bit_longer_without_loading_all_data_first ..")
98
99     def get_chunk(self, is_train=True):
100         data = self.training_data if is_train else self.testing_data
101         labels = self.training_labels if is_train else self.testing_labels
102
103         i = randint(0, len(data))
104         x = data[i] # should be a list at this point
105         label = labels[i]
106
107         if not load_all_to_memory: # means that data is a path instead
108             fs, x = wavfile.read(x[0])
109             label = get_label(label[0])
110         start = randint(0, len(label) - 1)

```

```

111     return x.T[:, start * window_size:start * window_size + window_size], label[start]
112
113 def get_complete_file_chunk( self , is_train =True):
114     data = self.training_data if is_train else self.testing_data
115     labels = self.training_labels if is_train else self.testing_labels
116
117     i = randint(0, len(data))
118     x = data[i] # should be a list at this point
119     label = labels[i]
120
121     if not load_all_to_memory: # means that data is a path instead
122         fs, x = wavfile.read(x[0])
123         label = get_label(label[0])
124     start = randint(0, len(label) - 1)
125     return x.T[:, start * window_size:start * window_size + window_size], label[start]
126
127 def load_all_data_to_memory(self):
128     paths = array(self.training_data, copy=True)
129     self.training_data = []
130     for path in paths:
131         fs, x = wavfile.read(path[0])
132         self.training_data.append(x)
133
134     paths = array(self.testing_data, copy=True)
135     self.testing_data = []
136     for path in paths:
137         fs, x = wavfile.read(path[0])
138         self.testing_data.append(x)
139
140     paths = array(self.training_labels, copy=True)
141     self.training_labels = []
142     for path in paths:
143         self.training_labels.append(get_label(path[0]))
144
145     paths = array(self.testing_labels, copy=True)
146     self.testing_labels = []
147     for path in paths:
148         self.testing_labels.append(get_label(path[0]))
149
150 def split_train_test ( self ):
151     shuf = vstack((self.data_paths, self.label_paths)).T
152     shuffle (shuf)
153
154     self.imax = i_max = int(len(self.data_paths[0]) * percent_train) + 1
155     self.training_data = shuf[0:i_max, [0]]
156     self.training_labels = shuf[0:i_max, [1]]
157     self.testing_data = shuf[i_max - 1:-1, [0]]
158     self.testing_labels = shuf[i_max - 1:-1, [1]]
159     # At this point the training and testing sets contain a list of paths to data and labels.
160
161 def train( self ):
162     try:
163         if load_path is not None:
164             print('Loading_Diarization')
165             diarization_method.load(load_path)
166
167
168     X = array([], dtype=float).reshape(0,window_size)

```



```

169         Y = array([], dtype=int).reshape(0,1)
170         for i in range(epochs):
171             for j in range(batch_size):
172                 x, label = self.get_chunk()
173                 X = vstack((X,x))
174                 Y = vstack((Y,label))
175                 diarization_method.train_on_data(X, Y)
176
177         print( 'Finished Training ... ' )
178         print( 'Testing ... ' )
179
180         # Testing section
181         l = len(self.testing_data)
182         for i in range(l):
183             x, label = self.get_chunk(is_train=False)
184             X = vstack((X,x))
185             Y = vstack((Y,label))
186             num_error = diarization_method.get_train_error(X, Y)
187             num_train = Y.shape[0]
188             training_accuracy = (num_train-num_error)/num_train*100
189             print( "%d/%d Mistakes. Training Accuracy: %.2f%%"%(int(num_error),num_train,
190                 training_accuracy))
191
192         except Exception as ex:
193             print( "ERROR! And exception Occurred!" )
194             print( ex )
195         except KeyboardInterrupt:
196             print( "\n Interrupted.." )
197         finally :
198             if save_path is not None:
199                 print( ' Saving Diarization to %s...'%(save_path))
200                 diarization_method.save(save_path)
201
202 class Predictor:
203     def __init__( self ):
204         self.data_paths = run_audio_file_path
205         self.label_paths = None
206         self.data = []
207         self.labels = []
208         self.imax = -1
209         self.textgrid = None
210
211         if text_grids_path is not None:
212             self.labels = get_label( text_grid_load_file_path )[1:-1]
213
214         fs, x = wavfile.read( self.data_paths )
215         X0 = self.create_data_chunks(x[:,0], window_size,window_size)
216         X1 = self.create_data_chunks(x[:,1], window_size,window_size)
217         self.data = hstack((X0,X1))
218
219     def create_data_chunks( self ,x,window_size,step_size ):
220         window_size = int(window_size)
221         step_size = int( step_size )
222
223         # drop elements outside window
224         if int(len(x)%(step_size)) != 0:
225             x = x[: -int(len(x)%(step_size))]

```

```

226
227     # ending
228     data_nb = int(ceil((len(x) - window_size)/step_size))
229
230     start=0
231     end=window_size
232     data = zeros([data_nb,window_size])
233
234     for i in range(data_nb):
235         data[i,:] = x[start:end]
236         start=start+step_size
237         end=end+step_size
238
239     return data
240
241 def run(self):
242     try:
243         if load_path is not None:
244             print('Loading_Diarization')
245             diarization_method.load(load_path)
246
247             # diarization_method.run_on_data(self.data_paths)
248             num_error = diarization_method.get_train_error(self.data, self.labels)
249             num_train = self.labels.shape[0]
250             training_accuracy = (num_train-num_error)/num_train*100
251             print("%d/%d_Mistakes._Training_Accuracy:_.2f%%"%(int(num_error),num_train,
252                 training_accuracy))
253
254         except Exception as ex:
255             print("ERROR!_and_Exception_Occurred!")
256             print(ex)
257         except KeyboardInterrupt:
258             print("\n_Interrupted..")
259         finally :
260             print("Done")
261
262 if __name__ == '__main__':
263     # TODO: command line argument parsing. Use are parse
264     if training_flag is True:
265         trainer = Trainer()
266         trainer.train()
267     else:
268         predictor = Predictor()
269         predictor.run()

```

A.2 DiarizationBaseClass.py

```

1 from abc import ABC, abstractmethod
2 import tensorflow as tf
3
4
5 class DiarizationBaseClass(ABC):
6     """
7     The DiarizationBaseClass provides the interfaces between custom diarization
8     methods and the wrapper code. It provides a tensorflow session available as
9     'self.sess' that is usable within the 'train_on_data', 'get_train_error', and

```

```

10 'run_on_data' methods.
11
12 You may add other variables needed within your own implementation. Not here
13 please. You can pass parameters into your custom diarization method via the
14 'params' init parameter. This can be an array, dictionary, or object. In
15 other words, that is the only parameter you'll ever need. You can pack
16 everything into it.
17 """
18
19 def __init__(self, params=None):
20     """
21     Initialization for the base diarization class.
22     :param params: the tuple, dictionary, or object used to initialize the
23                   diarization method.
24     """
25     self.tf_initializer = tf.global_variables_initializer()
26     self.sess = None
27     self.init_diarization_method(params)
28     pass
29
30 @abstractmethod
31 def init_diarization_method(self, params):
32     """
33     This method is called once to initialize the speech diarization algorithm.
34     :param params: the tuple, dictionary, or object used to initialize the
35                   diarization method.
36     """
37     pass
38
39 @abstractmethod
40 def train_on_data(self, data, label):
41     """
42     This method is meant to be called once per iteration during the training cycle.
43
44     The tensorflow session is available as 'self.sess'. tensorflow session initialization
45     has already been taken care of with the 'tf.global_variables_initializer()'
46
47     :param data: This is the single set of data to be used by a single training
48                 operation. This will change with each call of 'train_on_data'.
49     :param label: This is the textgrid label for this data sample.
50     """
51     pass
52
53 @abstractmethod
54 def run_on_data(self, data):
55     """
56     This method will be used to run the diarization method on a set of data and
57     returns the TextGrid result.
58     :param data: the audiofile to run the diarization on.
59     :return: the TextGrid object.
60     """
61     pass
62
63 @abstractmethod
64 def get_train_error(self, test_data, test_label):
65     """
66     This method will be called at the end of each epoch to get the error of the
67     diarization method.

```

```

68         :param test_data: the data to use to evaluate the accuracy of the model.
69         :param test_label: the label to use to evaluate the accuracy of the model.
70         :return: The error of the diarization method.
71         """
72
73     @abstractmethod
74     def load(self, path):
75         """
76         This method will be called when we want to load the diarization method from disk
77         :param path: the path to the storage file that is to be loaded.
78         :return: True on success, False on failure
79         """
80         pass
81
82     @abstractmethod
83     def save(self, path):
84         """
85         This method will be called when we want to save the diarization method to disk to
86         be recalled later .
87         :param path: the path to where the storage file should be put.
88         :return: True on success, False on failure
89         """

```

A.3 TFLearnNN.py

```

1  from diarization_methods.DiarizationBaseClass import DiarizationBaseClass
2  import tensorflow as tf
3  import tflearn
4  from tflearn.layers.core import input_data, dropout, fully_connected
5  from tflearn.layers.conv import conv_2d, max_pool_2d
6  from tflearn.layers.normalization import local_response_normalization
7  from tflearn.layers.estimator import regression
8  import time
9  import numpy as np
10
11 from numpy import array, zeros
12
13
14 class TFLearnNN(DiarizationBaseClass):
15     def run_on_data(self, data):
16         # create fft
17         if len(test_data) != len(test_label):
18             test_data = test_data.reshape(len(test_label), -1)
19
20         Xfft = self._get_fft(test_data)
21         X = np.hstack((test_data, Xfft))
22
23         # create one-hot label
24         Y = self.label_to_one_hot(test_label, self.num_classes)
25
26         yhat = self.model.predict(X)
27         yhat = np.argmax(yhat, axis=1)
28
29         return self._validate_results(yhat, test_label)
30
31     def label_to_one_hot(self, label, num_classes):
32         """

```

```

33     converts
34     into
35     :param label: [0, 3, 2, 1, ...]
36     :return:      [[1, 0, 0, 0, ...],
37                   [0, 0, 0, 1, ...],
38                   [0, 0, 1, 0, ...],
39                   [0, 1, 0, 0, ...]]
40     """
41     y = np.eye(self.num_classes)[label.astype(int)].reshape(-1,num_classes)
42     return y
43
44 def init_diarization_method( self , params):
45     self.window_size = params['window_size']
46     self.nfft         = params['nfft']
47     self.num_channels = params['num_channels']
48     self.num_classes  = params['num_classes']
49     self.total_width  = (self.window_size + self.nfft)*self.num_channels
50
51     input_layer = tflearn.input_data(shape=[None, self.total_width], name='input')
52     dense1 = tflearn.fully_connected(input_layer, 300, activation='relu')
53
54     dense2 = tflearn.fully_connected(dense1, 200, activation='relu')
55     dropout2 = tflearn.dropout(dense2, 0.8)
56
57     dense3 = tflearn.fully_connected(dropout2, 200, activation='relu')
58     dropout3 = tflearn.dropout(dense3, 0.8)
59
60     dense4 = tflearn.fully_connected(dropout3, 500, activation='relu')
61     dropout4 = tflearn.dropout(dense4, 0.8)
62
63     softmax = tflearn.fully_connected(dropout4, 4, activation='softmax')
64
65     self.net = tflearn.regression(softmax, loss='categorical_crossentropy', name='target')
66     self.model = tflearn.DNN(self.net, tensorboard_verbose=2)
67
68 def train_on_data( self , data, label):
69     if len(data) != len(label):
70         data = data.reshape(len(label),-1)
71         test_data = test_data.reshape(len( test_label ),-1)
72
73     # create fft
74     Xfft = self._get_fft (data)
75     X = np.hstack((data,Xfft))
76
77     # create one-hot label
78     Y = self.label_to_one_hot (label, self.num_classes)
79
80     self.model.fit(X, Y, n_epoch=1, show_metric=True, \
81                   validation_set=(X, Y),
82                   snapshot_epoch=False,run_id='TfLearnNNk')
83
84 def get_train_error ( self , test_data, test_label):
85     # create fft
86     if len(test_data) != len( test_label ):
87         test_data = test_data.reshape(len( test_label ),-1)
88
89     Xfft = self._get_fft (test_data)
90     X = np.hstack((test_data,Xfft))

```

```

91
92     # create one-hot label
93     Y = self.label_to_one_hot ( test_label , self . num_classes)
94
95     yhat = self.model.predict(X)
96     yhat = np.argmax(yhat, axis=1)
97
98     return self . __validate_results (yhat, test_label )
99
100 def load( self , path):
101     self . model.load(path)
102
103 def save( self , path):
104     self . model.save(path)
105
106 def __validate_results ( self , yhat, y):
107     num_error = 0
108     for i in range(len(yhat)):
109         if yhat[i] != y[i]:
110             num_error += 1
111     return num_error
112
113 def __get_fft ( self , data):
114     if self . num_channels is 2:
115         Xfft0 = abs(np.fft . fft (data[:, self . window_size], self . nfft ))
116         Xfft1 = abs(np.fft . fft (data[:, self . window_size:], self . nfft ))
117         Xfft = np.hstack((Xfft0,Xfft1))
118     else :
119         Xfft = abs(np.fft . fft (data[:, self . window_size:], self . nfft ))
120     return Xfft

```

A.4 plot.py

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3 from scipy.io import wavfile
4 import scipy.fftpack as fftpack
5 from scipy import signal
6 import sounddevice as sd
7 import matplotlib.animation as animation
8 from matplotlib import rcParams
9
10 rcParams.update({'figure.autolayout': True})
11
12
13 def track_audio(fs , start=0, end="end"):
14     Y_MIN = -2
15     Y_MAX = 2
16
17     x = np.arange(start, end + 1, 0.01);
18
19     def update_line(num, line):
20         i = x[num]
21         line.set_data([i, i], [Y_MIN, Y_MAX])
22     return line ,
23
24     l, _ = plt.plot(start , -1, end, 1, linewidth=2, color='red')

```

```

25 line_anim = animation.FuncAnimation(fig, update_line, len(x), fargs=(1,), interval=1 / fs, blit
    =True, repeat=False)
26 plt.show()
27
28
29 def plot_spectrogram(x, fs, start=0, end="end"):
30     if end is "end":
31         Pxx, freqs, bins, im = plt.specgram(x[start * fs:-1],
32                                             NFFT=512, Fs=fs, noverlap=100, cmap=plt.cm.gist_heat
33                                             )
34         end = max(bins)
35     else:
36         Pxx, freqs, bins, im = plt.specgram(x[start * fs:start * fs + end * fs],
37                                             NFFT=512, Fs=fs, noverlap=100, cmap=plt.cm.gist_heat
38                                             )
39         if end > max(bins):
40             end = max(bins)
41
42     plt.ylim(0, max(freqs))
43     plt.xlim([start, end])
44     plt.ylabel("Frequency_(Hz)")
45     plt.xlabel("Time_(Sec)")
46     # plt.colorbar()
47
48 def play_audio(data, fs, start=0, end="end", blocking=True):
49     if end is "end":
50         sd.play(data[int(start * fs):-1], fs, blocking=blocking)
51     else:
52         sd.play(data[start * fs:int(start * fs + end * fs)], fs, blocking=blocking)
53
54
55 def plot_audio(x, fs, start=0, end="end"):
56     if end is "end":
57         xaxis = np.linspace(start, len(x) / fs, num=len(x[start * fs:]))
58         plt.plot(xaxis, x[start * fs:] / float(max(x)), linewidth=0.25)
59         end = len(x) / float(fs)
60
61     else:
62         if end > len(x) / fs:
63             end = len(x) / fs
64         xaxis = np.linspace(start, end, num=len(x[int(start * fs):int(end * fs)]))
65         plt.plot(xaxis, x[int(start * fs):int(end * fs)] / float(max(x)), color='#3030e0',
66                 linewidth=0.15)
67
68     plt.ylim(-1.2, 1.2)
69     plt.xlim([start, end])
70     plt.ylabel("Magnitude")
71     plt.xlabel("Time_(Sec)")
72
73 def plot_bounds_lines(changes, marks, start=0, end="end"):
74     x = changes
75     if end is "end":
76         last = x.shape[0] + 1
77     else:
78         last = np.argmax(x > end)

```

```

79
80 data_range = range(np.argmin(x < start), last)
81
82 if max(marks) > 2:
83     plt.plot((-1, -1), (-1, -1), 'b-', linewidth=2, label="Speaker_1")
84     plt.plot((-1, -1), (-1, -1), 'r-', linewidth=2, label="Speaker_2")
85     plt.plot((-1, -1), (-1, -1), 'g-', linewidth=2, label="Both")
86 else:
87     plt.plot((-1, -1), (-1, -1), 'b-', linewidth=2, label="Speech")
88
89 plt.legend(loc=2, fancybox=True, framealpha=0.8, prop={'size': 9})
90
91 for j, mark in enumerate(marks):
92     if mark == 1:
93         plt.plot((x[j], x[j + 1]), (-1.05, -1.05), 'b-', linewidth=2)
94         plt.plot((x[j + 1], x[j + 1]), (-1.03, -1.07), 'k-', linewidth=2, zorder=10)
95
96     elif mark == 2:
97         plt.plot((x[j], x[j + 1]), (-1.05, -1.05), 'r-', linewidth=2)
98         plt.plot((x[j + 1], x[j + 1]), (-1.03, -1.07), 'k-', linewidth=2, zorder=10)
99
100    elif mark == 3:
101        plt.plot((x[j], x[j + 1]), (-1.05, -1.05), 'g-', linewidth=2)
102        plt.plot((x[j + 1], x[j + 1]), (-1.03, -1.07), 'k-', linewidth=2, zorder=10)
103
104    else:
105        pass
106
107    plt.ylim(-1.2, 1.2)
108    plt.xlim([start, end])
109    plt.xlabel("Time_(Sec)")
110
111
112 def plot_bounds_fill(changes, marks, start=0, end="end"):
113     x = changes
114     if end is "end":
115         last = x.shape[0] + 1
116     else:
117         last = np.argmax(x >= end) + 1
118
119     data_range = range(np.argmin(x < start), last)
120
121     # create legend information
122     # plt.axvspan(0, 0, alpha=0.5, color='b', label="\N\" Human")
123     # plt.axvspan(0, 0, alpha=0.5, color='r', label="\S\" Human")
124     # plt.axvspan(0, 0, alpha=0.5, color='g', label="\S\" Human")
125     plt.legend(loc=2, fancybox=True, framealpha=0.8, prop={'size': 9})
126
127     for j, mark in enumerate(marks):
128         if mark == 1:
129             plt.axvspan(x[j], x[j + 1], alpha=0.2, color='b')
130
131         elif mark == 2:
132             plt.axvspan(x[j], x[j + 1], alpha=0.2, color='r')
133
134         elif mark == 3:
135             plt.axvspan(x[j], x[j + 1], alpha=0.2, color='g')
136

```



```

137     else :
138         pass
139
140
141 def plot_fft () :
142     N = 600 # sample points
143     T = 1 / 800.0 # sample spacing
144     x = np.linspace(0, N * T, N)
145     y = np.sin(50.0 * 2.0 * np.pi * x) + 0.5 * np.sin(80.0 * 2.0 * np.pi * x)
146     yf = fftpack.fft(y)
147     xf = np.linspace(0.0, 1.0 / (2.0 * T), N / 2)
148     fig, ax = plt.subplots()
149     ax.plot(xf, 2.0 / N * np.abs(yf[:N // 2]))
150
151
152 if __name__ == '__main__':
153     main()

```

A.5 textgrid.py

An existing TextGrid parser was used from <https://github.com/kylebgorman/textgrid/blob/master/textgrid/textgrid.py>. Three additional functions were added to assist in neural network training:

```

1 def get_bin_changes(x, Fs):
2     FsTimeChanges = [0]
3     FsChangeMarks = [x[0]]
4     for i in range(len(x)):
5         if x[i] != x[i - 1]:
6             FsTimeChanges.append(i)
7             FsChangeMarks.append(x[i])
8
9     FsTimeChanges.append(i)
10    FsTimeChanges = np.array(FsTimeChanges)*1/float(Fs)
11    FsChangeMarks = np.array(FsChangeMarks)
12    return FsTimeChanges, FsChangeMarks

```

```

1 class IntervalTier(object):
2     ...
3     def tier_to_array (self, Fs, maxTime):
4         array = np.zeros(int(np.ceil (Fs * maxTime)))
5         current_interval = 0
6         for i in range(array.shape[0]):
7             if i/float (Fs) < self[ current_interval ].maxTime:
8                 if self[ current_interval ].mark == 'N':
9                     array[i] = 1
10                else :
11                    current_interval += 1
12
13    self.FsArray = array
14    self.FsTimeChanges, self.FsChangeMarks = get_bin_changes(array, Fs)
15    ...

```

```

1 class TextGrid(object):
2     ...
3
4     def combine_grids(self):

```

```

5 self.FsArrayCombined = self.tiers[0].FsArray + self.tiers[1].FsArray * 2
6 self.FsTimeChangesCombined, self.FsChangeMarksCombined = get_bin_changes(self.
    FsArrayCombined, self.Fs)
7 ...

1 class TextGrid(object):
2 ...
3 def read(self, f, round_digits=DEFAULT_TEXTGRID_PRECISION, Fs=None):
4     """
5     Read the tiers contained in the Praat-formatted TextGrid file
6     indicated by string f. Times are rounded to the specified precision.
7     """
8     self.Fs = Fs
9     encoding = detectEncoding(f)
10    with codecs.open(f, 'r', encoding=encoding) as source:
11        source.readline() # header junk
12        source.readline() # header junk
13        source.readline() # header junk
14
15        self.minTime = round(float(source.readline().split()[2]), round_digits)
16        self.maxTime = round(float(source.readline().split()[2]), round_digits)
17        source.readline() # more header junk
18        m = int(source.readline().rstrip().split()[2]) # will be self.n
19        source.readline()
20        for i in range(m): # loop over grids
21            source.readline()
22            if source.readline().rstrip().split()[2] == "IntervalTier":
23                inam = source.readline().rstrip().split(' _= _')[1].strip(' ')
24                imin = round(float(source.readline().rstrip().split()[2]), round_digits)
25                imax = round(float(source.readline().rstrip().split()[2]), round_digits)
26                itie = IntervalTier(inam)
27                for j in range(int(source.readline().rstrip().split()[3])):
28                    source.readline().rstrip().split() # header junk
29                    jmin = round(float(source.readline().rstrip().split()[2]), round_digits)
30                    jmax = round(float(source.readline().rstrip().split()[2]), round_digits)
31                    jmrk = _getMark(source)
32                    if jmin < jmax: # non-null
33                        itie.addInterval(Interval(jmin, jmax, jmrk))
34                self.append(itie)
35            else: # pointTier
36                inam = source.readline().rstrip().split(' _= _')[1].strip(' ')
37                imin = round(float(source.readline().rstrip().split()[2]), round_digits)
38                imax = round(float(source.readline().rstrip().split()[2]), round_digits)
39                itie = PointTier(inam)
40                n = int(source.readline().rstrip().split()[3])
41                for j in range(n):
42                    source.readline().rstrip() # header junk
43                    jtim = round(float(source.readline().rstrip().split()[2]), round_digits)
44                    jmrk = _getMark(source)
45                    itie.addPoint(Point(jtim, jmrk))
46                self.append(itie)
47
48    if Fs != None:
49        for tier in self.tiers:
50            tier.tier_to_array(Fs, self.maxTime)
51    if len(self.tiers) == 2:
52        self.combine_grids()
53    ...

```